

✓ LoRA MLP + Attention (Low Rank)

This notebook mirrors the full-run LoRA pipeline and focuses on low-rank target-module ablations for attention-only vs attention+MLP adapters.

```
from pathlib import Path
import sys
import pandas as pd
import torch
from IPython.display import display
```

✓ Colab setup and project root

```
try:
    from google.colab import drive
    IN_COLAB = True
except ImportError:
    IN_COLAB = False

if IN_COLAB:
    drive.mount('/content/drive')
    CANDIDATE_ROOT = Path('/content/drive/MyDrive/DL_Final_Project')
    PROJECT_ROOT = CANDIDATE_ROOT if CANDIDATE_ROOT.exists() else Path.
else:
    cwd = Path.cwd().resolve()
    if cwd.name == 'notebooks':
        PROJECT_ROOT = cwd.parent
    elif cwd.name == 'LoRA_study' and cwd.parent.name == 'notebooks':
        PROJECT_ROOT = cwd.parent.parent
    else:
        PROJECT_ROOT = cwd

print('IN_COLAB:', IN_COLAB)
print('PROJECT_ROOT:', PROJECT_ROOT)
if str(PROJECT_ROOT) not in sys.path:
    sys.path.append(str(PROJECT_ROOT))
```

```
Drive already mounted at /content/drive; to attempt to forcibly remoun
IN_COLAB: True
PROJECT_ROOT: /content/drive/MyDrive/DL_Final_Project
```

```

import importlib.metadata as md
import importlib.util
import subprocess

def _parse_version_tuple(version_str: str) -> tuple[int, ...]:
    parts = []
    for token in version_str.split('.'):
        digits = ''.join(ch for ch in token if ch.isdigit())
        parts.append(int(digits) if digits else 0)
    return tuple(parts)

required_torchao = (0, 16, 0)
if importlib.util.find_spec('torchao') is not None:
    v_str = md.version('torchao')
    print(f'Found torchao=={v_str}')
    if _parse_version_tuple(v_str) < required_torchao:
        print('Uninstalling incompatible torchao...')
        subprocess.check_call([sys.executable, '-m', 'pip', 'uninstall',
                                'torchao'])
        raise SystemExit('Removed incompatible torchao. Restart runtime')
    print('torchao is compatible.')
else:
    print('torchao is not installed (this is fine).')

```

```
torchao is not installed (this is fine).
```

```

from src.gen_lora_study_utils import (
    set_seed,
    load_split,
    apply_sanity_subset,
    cap_validation_rows,
)
from src.lora_best_candidates_full_run_utils import build_configs_from
from src.lora_tuning_block_utils import count_trainable_params_for_config

```

✓ Runtime controls (same style as full-run notebook)

```

save = True

# For the sweep, keep this False so you do not run test inference for

```

```

# After choosing the best validation config, rerun only that config w
generate_submission = True #True

# Empty means "all ablations" (same behavior as submission_ablation_id
submission_ablation_ids = []
save_val_predictions = True
val_prediction_ablation_ids = []#['score_lora_sft', 'score_lora_option

sanity_check = False
n = 313 # ignored when sanity_check=False

seed = 42
max_val_examples = 256 #None
top_k = 5
batch_size = 4
clear_cuda_between_runs = True

# For option_scoring-only runs, loss is a smoother early-stopping sign
# For mixed option_scoring + SFT runs, use 'accuracy' instead; see no
early_stopping = None
epochs_per_val_check = None
early_stopping_metric = 'loss'
early_stopping_val_examples = 128 #None

DATA_DIR = PROJECT_ROOT / 'data'
IMAGE_DIR = DATA_DIR / 'images' / 'images'
OUT_DIR = PROJECT_ROOT / 'outputs' / '03_dora_scoring_long_run'
ADAPTER_DIR = OUT_DIR / 'adapters'
MODEL_ID = 'HuggingFaceTB/SmolVLM-500M-Instruct'
DEVICE = 'cuda' if torch.cuda.is_available() else 'cpu'

torch_dtype = (
    'bfloat16'
    if torch.cuda.is_available() and torch.cuda.is_bf16_supported()
    else ('float16' if torch.cuda.is_available() else 'auto')
)

print('CUDA available:', torch.cuda.is_available())
if torch.cuda.is_available():
    print('GPU:', torch.cuda.get_device_name(0))
print({
    'save': save,
    'generate_submission': generate_submission,
    'submission_ablation_ids': submission_ablation_ids,
    'save_val_predictions': save_val_predictions,

```

```

        'val_prediction_ablation_ids': val_prediction_ablation_ids,
        'sanity_check': sanity_check,
        'n': n,
        'device': DEVICE,
    })
    set_seed(seed)

```

```

CUDA available: True
GPU: NVIDIA A100-SXM4-80GB
{'save': True, 'generate_submission': True, 'submission_ablation_ids':

```

```

train_df = load_split(DATA_DIR, 'train')
val_df = load_split(DATA_DIR, 'val')
test_df = load_split(DATA_DIR, 'test')

train_df = apply_sanity_subset(train_df, sanity_check, n, seed)
val_df = apply_sanity_subset(val_df, sanity_check, n, seed)
test_df = apply_sanity_subset(test_df, sanity_check, 1008, seed)
val_df = cap_validation_rows(val_df, max_val_examples)

print('train rows used:', len(train_df))
print('validation rows used:', len(val_df))
print('test rows used:', len(test_df))
print('IMAGE_DIR exists:', IMAGE_DIR.exists())

```

```

train rows used: 3109
validation rows used: 256
test rows used: 1008
IMAGE_DIR exists: True

```

✓ Candidate ablation spaces

The key differences are low-rank settings and target modules. Input/data ablations are toggled directly in each candidate config.

```

# Single best-attempt config after the SFT masking / generation / opti

input_ablation_1 = {
    # The strategy note says the original images are small and recomm
    # Use 512 for this run. If you later find many images are non-squ
    # and control processor resolution instead.
    'image_size': 200,

```

```

# Keep augmentation light. These images likely contain diagrams/tables
'enable_image_augmentation': True,
'brightness_jitter': 0.03,
'slight_rotation_deg': 0.5,

# Use metadata because subject/grade/topic may help route the request
'include_metadata_in_prompt': True,
'metadata_fields': ['subject', 'grade', 'topic'],
'metadata_header': 'Metadata',

# With output_format='letter_only', this is mostly harmless/unused
# but keep it simple in case other code references it.
'answer_prefix_text': 'Answer:',
'question_label': 'Question:',
'choices_label': 'Choices:',
'instruction_prefix': '',
}
input_ablation_2 = {
    # The strategy note says the original images are small and recommended
    # Use 512 for this run. If you later find many images are non-square
    # and control processor resolution instead.
    'image_size': 384,

    # Keep augmentation light. These images likely contain diagrams/tables
    'enable_image_augmentation': True,
    'brightness_jitter': 0.03,
    'slight_rotation_deg': 0.5,

    # Use metadata because subject/grade/topic may help route the request
    'include_metadata_in_prompt': True,
    'metadata_fields': ['subject', 'grade', 'topic'],
    'metadata_header': 'Metadata',

    # With output_format='letter_only', this is mostly harmless/unused
    # but keep it simple in case other code references it.
    'answer_prefix_text': 'Answer:',
    'question_label': 'Question:',
    'choices_label': 'Choices:',
    'instruction_prefix': '',
}
prompt_structure = 'question_first'
context_mode = 'hint_only'
choice_format = 'letter_dot'

```

```

# Important: use letter-only with option scoring.
# This makes each candidate just "A", "B", "C", etc.

output_format = 'letter_only'

# Option scoring does not rely on free generation, but keep these same
max_new_tokens = 4
decoding_strategy = 'greedy'
num_beams = 1
parse_rule = 'strict_first_letter'

# Option scoring is slower than SFT because it scores every candidate
# Start with 12 rather than 25; early stopping can still stop sooner.
epochs = 5

ablation_space = {
    'dora': {
        # Main change: use multiple-choice candidate scoring rather than
        'training_objective': 'option_scoring',

        'prompt_structure': prompt_structure,
        'context_mode': context_mode,
        'choice_format': choice_format,
        'output_format': output_format,
        'max_new_tokens': max_new_tokens,
        'decoding_strategy': decoding_strategy,
        'num_beams': num_beams,
        'parse_rule': parse_rule,

        # DoRA + attention/MLP targets.
        # r=4 should stay safely below your 5M trainable-parameter cap
        # attention behavior and internal MLP representations.
        'lora_r': 4,
        'lora_alpha': 8,
        'lora_dropout': 0.03,
        'target_modules': [
            'q_proj', 'k_proj', 'v_proj', 'o_proj',
        ],
        'lora_bias': 'none',
        'lora_task_type': 'CAUSAL_LM',
        'use_dora': False,
        'use_rslora': True,

        'epochs': epochs,
    },
}

```

```

# Your current option-scoring trainer is effectively row-by-row
# may not be used unless you patched it in. Keep this at 1.
'gradient_accumulation_steps': 1,

# DoRA usually benefits from a slightly lower LR than vanilla
'lr': 5e-5,
'weight_decay': 0.01,
'max_grad_norm': 1.0,

# These only matter if you added scheduler support.
# If not, they will remain harmless config fields.
'warmup_ratio': 0.05,
'scheduler_type': 'cosine',

**input_ablation_2,
},
}

all_configs = build_configs_from_named_space(ablation_space, 'score')
display(pd.DataFrame(all_configs))

```

	training_objective	prompt_structure	context_mode	choice_format	on
0	option_scoring	question_first	hint_only	letter_dot	

1 rows x 37 columns

```

fixed_context = {
    'MODEL_ID': MODEL_ID,
    'DEVICE': DEVICE,
    'torch_dtype': torch_dtype,
    'IMAGE_DIR': IMAGE_DIR,
    'train_df': train_df,
    'val_df': val_df,
    'test_df': test_df,
    'batch_size': batch_size,
    'clear_cuda_between_runs': clear_cuda_between_runs,
    'early_stopping': early_stopping,
    'epochs_per_val_check': epochs_per_val_check,
    'early_stopping_metric': early_stopping_metric,
    'early_stopping_val_examples': early_stopping_val_examples,
}

```



```

# Save adapters once during val training so test submissions can
'save_trained_adapters': bool(generate_submission),
'ADAPTER_DIR': ADAPTER_DIR,
}

_param_rows = []
for cfg in all_configs:
    counts = count_trainable_params_for_config(cfg, {'MODEL_ID': MODEL_ID})
    _param_rows.append({
        'ablation_id': cfg['ablation_id'],
        'trainable_params': counts['trainable'],
        'total_params': counts['total'],
        'trainable_pct': round(float(counts['trainable_pct']), 4),
    })
display(pd.DataFrame(_param_rows).sort_values('ablation_id').reset_index())

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:
The secret `HF\_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings page.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access most features.

warnings.warn(
Warning: You are sending unauthenticated requests to the HF Hub. Please
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated
config.json: 7.34k/? [00:00<00:00, 782kB/s]

`torch\_dtype` is deprecated! Use `dtype` instead!

model.safetensors: 100% 1.02G/1.02G [00:04<00:00, 722MB/s]

Loading weights: 100% 489/489 [00:00<00:00, 1125.97it/s, Materializing param=model.vision_model.encoder.layers.10.self_attn.q_proj.weight]

generation_config.json: 100% 136/136 [00:00<00:00, 12.9kB/s]

	ablation_id	trainable_params	total_params	trainable_pct
--	--------------------	-------------------------	---------------------	----------------------

0	score_dora	1040384	508522688	0.2046
----------	------------	---------	-----------	--------

```

results = []
val_preds = []
histories = []
submission_frames = []

```

```

config_by_id = {cfg['ablation_id']: cfg for cfg in all_configs}
all_ablation_ids = set(config_by_id.keys())

```

```

submission_requested_ids = [str(x) for x in (submission_ablation_ids (
submission_selected_ids = set(submission_requested_ids) if submission
submission_missing_ids = sorted(set(submission_requested_ids) - all_abl
if submission_missing_ids:
    print('Skipping unknown submission ablation ids:', submission_mis:

val_prediction_requested_ids = [str(x) for x in (val_prediction_ablat:
val_prediction_selected_ids = set(val_prediction_requested_ids) if val
val_prediction_missing_ids = sorted(set(val_prediction_requested_ids)
if val_prediction_missing_ids:
    print('Skipping unknown validation-prediction ablation ids:', val

for cfg in all_configs:
    ablation_id = cfg['ablation_id']
    summary, val_pred_df, history_df = run_lora_candidate_on_val(conf:
    results.append(summary)
    val_preds.append(val_pred_df)
    if len(history_df):
        histories.append(history_df)

# Generate submission immediately after training for selected abl:
if generate_submission and (ablation_id in submission_selected_ids:
    adapter_path = ADAPTER_DIR / ablation_id
    test_pred_df, submission_df = run_test_submission_from_saved:
        config=cfg,
        fixed_context=fixed_context,
        adapter_dir=adapter_path,
        save=save,
        out_dir=OUT_DIR,
        filename_prefix=f'mlp_attn_{ablation_id}_',
    )
    _sub = submission_df.copy()
    _sub['ablation_id'] = ablation_id
    submission_frames.append(_sub)
    print(f'Generated submission without retraining for: {ablation
    display(test_pred_df.head())
    display(submission_df.head())

if clear_cuda_between_runs and torch.cuda.is_available():
    torch.cuda.empty_cache()

results_df = pd.DataFrame(results).sort_values('accuracy', ascending=
val_predictions_df = pd.concat(val_preds, ignore_index=True) if val_p
history_df = pd.concat(histories, ignore_index=True) if histories else
display(results_df)

```

```

if submission_frames:
    combined_submission_preview = pd.concat(submission_frames, ignore
    display(combined_submission_preview.head())

if save and save_val_predictions:
    if val_prediction_requested_ids:
        selected_val_predictions_df = val_predictions_df[
            val_predictions_df['ablation_id'].isin(sorted(val_predict:
        ].reset_index(drop=True)
    else:
        selected_val_predictions_df = val_predictions_df.copy()

    if len(selected_val_predictions_df):
        OUT_DIR.mkdir(parents=True, exist_ok=True)
        val_pred_path = OUT_DIR / 'mlp_attn_val_predictions.csv'
        selected_val_predictions_df.to_csv(val_pred_path, index=False)
        print(
            'Saved validation predictions:',
            {
                'path': str(val_pred_path),
                'rows': int(len(selected_val_predictions_df)),
                'ablations': sorted(selected_val_predictions_df['abla
            },
        )
        display(selected_val_predictions_df.head())
    else:
        print('No validation predictions matched the selected ablation
elif save_val_predictions and not save:
    print('save_val_predictions=True but save=False; skipping validation

```

```

processor_config.json: 100% ██████████ 68.0/68.0 [00:00<00:00, 9.21kB/s]
chat_template.json: 100% ██████████ 429/429 [00:00<00:00, 54.8kB/s]
preprocessor_config.json: 100% ██████████ 486/486 [00:00<00:00, 63.4kB/s]
tokenizer_config.json: █████ 28.2k/? [00:00<00:00, 3.11MB/s]
tokenizer.json: █████ 3.55M/? [00:00<00:00, 123MB/s]
added_tokens.json: █████ 4.74k/? [00:00<00:00, 500kB/s]
special_tokens_map.json: █████ 1.07k/? [00:00<00:00, 113kB/s]
Loading weights: 100% ██████████ 489/489 [00:00<00:00, 1177.82it/s, Materializing param=model.vision_model

```

Score epoch 1/5: 100% 3109/3109 [2:45:36<00:00, 3.35s/it]

Score epoch 2/5: 100% 3109/3109 [2:09:46<00:00, 2.68s/it]

Score epoch 3/5: 100% 3109/3109 [2:10:05<00:00, 2.68s/it]

Score epoch 4/5: 100% 3109/3109 [2:09:53<00:00, 2.67s/it]

Score epoch 5/5: 100% 3109/3109 [2:09:57<00:00, 2.68s/it]

Val option scoring: 100% 256/256 [12:43<00:00, 2.82s/it]

Loading weights: 100% 489/489 [00:00<00:00, 1170.11it/s, Materializing param=model.vision_model

Val option scoring: 100% 1008/1008 [49:09<00:00, 2.33s/it]

Generated submission without retraining for: score_dora

	id	pred_answer	num_choices
0	test_01750	2	4
1	test_00128	0	5
2	test_02891	0	5
3	test_02425	1	5
4	test_00930	2	5

	id	answer
0	test_01750	2
1	test_00128	0
2	test_02891	0
3	test_02425	1
4	test_00930	2

	training_objective	prompt_structure	context_mode	choice_format
0	option_scoring	question_first	hint_only	letter_dot

1 rows x 46 columns

Start coding or [generate](#) with AI.

0	test_01750	2	score_dora
1	test_00128	0	score_dora
2	test_02891	0	score_dora

```
2 test_02891      0      score_dora
3 test_02425      1      score_dora
4 test_00930      2      score_dora
```

Saved validation predictions: {'path': '/content/drive/MyDrive/DL_Fin

	id	pred_answer	num_choices	gold_answer	is_correct	ablatic
0	val_00671	0	3	0	1	score
1	val_04111	0	3	1	0	score
2	val_02022	0	4	3	0	score
3	val_01237	0	5	0	1	score